

How Uber Builds A Cross Data Center Replication Platform on Apache Kafka

Hongliang Xu

Uber Streaming Data Team

Dec 07, 2018

Uber

Agenda

01 Apache Kafka at Uber

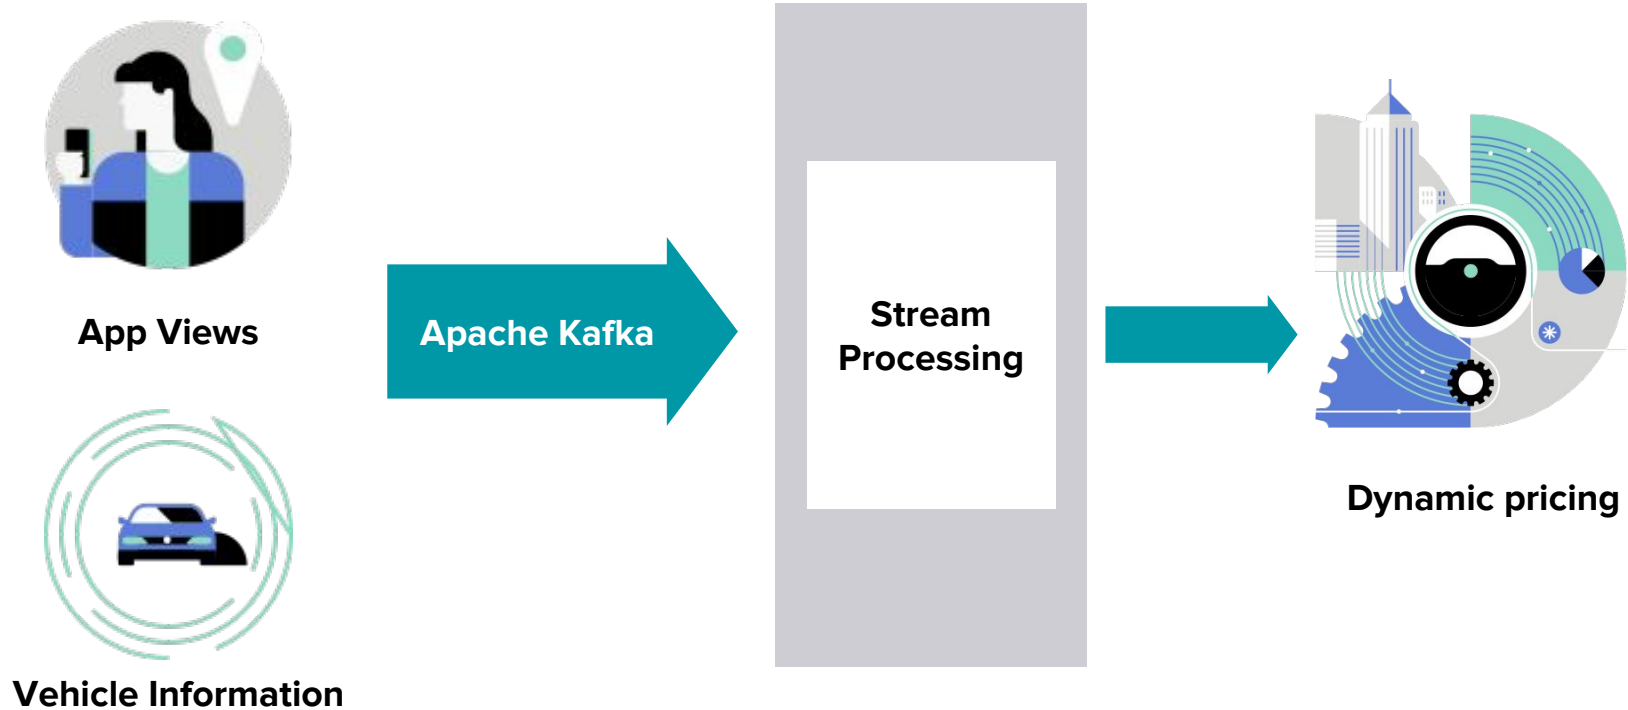
02 Apache Kafka pipeline & replication

03 uReplicator

04 Data loss detection

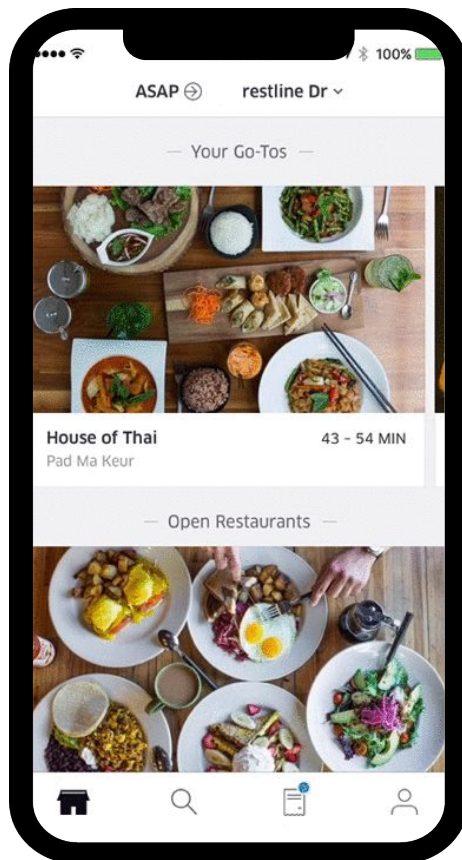
05 Q&A

Real-time Dynamic Pricing



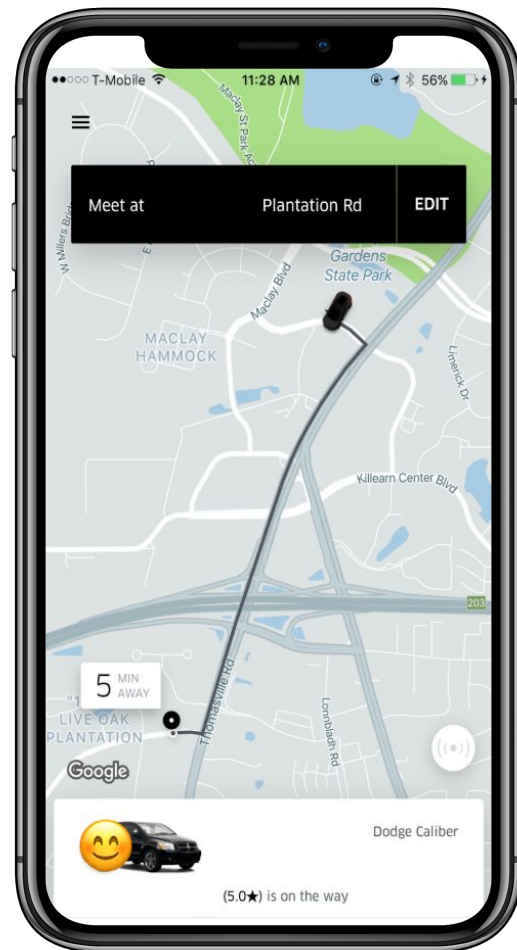
Uber Eats - Real-Time ETAs

Uber



A bunch more ...

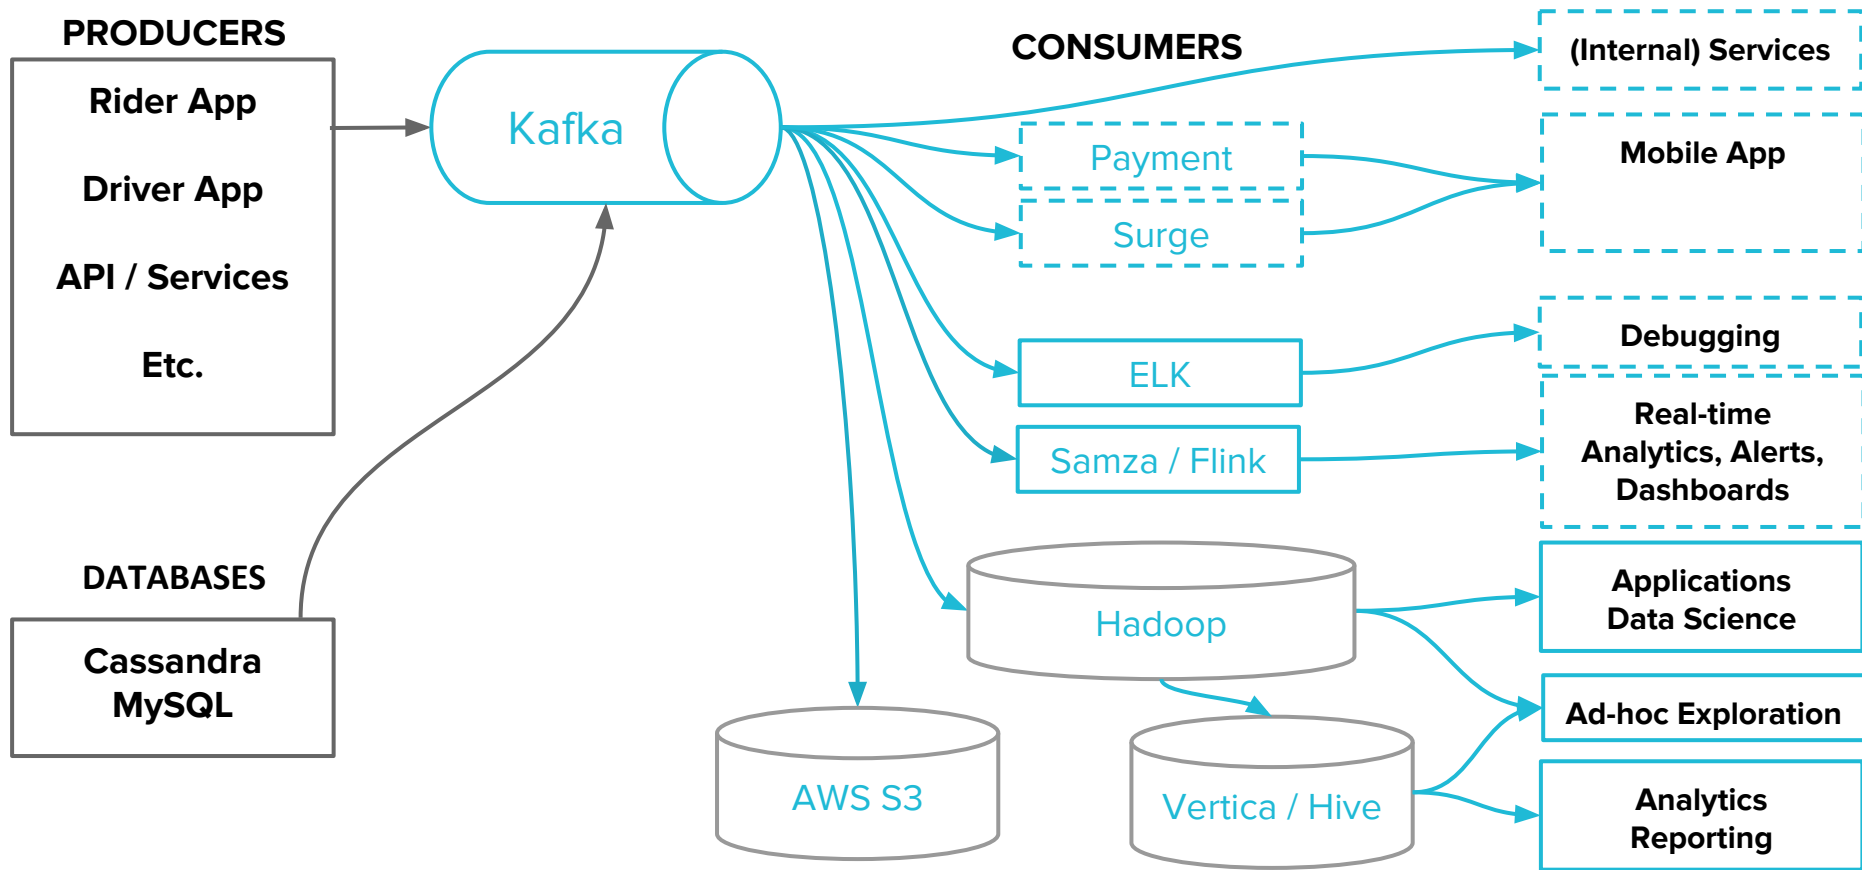
- **Fraud Detection**
- **Driver & Rider Sign-ups, etc.**



Apache Kafka - Use Cases

- General pub-sub, messaging queue
- Stream processing
 - AthenaX - self-service streaming analytics platform (Apache Samza & Apache Flink)
- Database changelog transport
 - Cassandra, MySQL, etc.
- Ingestion
 - HDFS, S3
- Logging

Data Infrastructure @ Uber



Scale

Uber

Trillions

Messages / Day

PBs

Data

Tens of
Thousands

Topics

excluding replication

Agenda

01 Apache Kafka at Uber

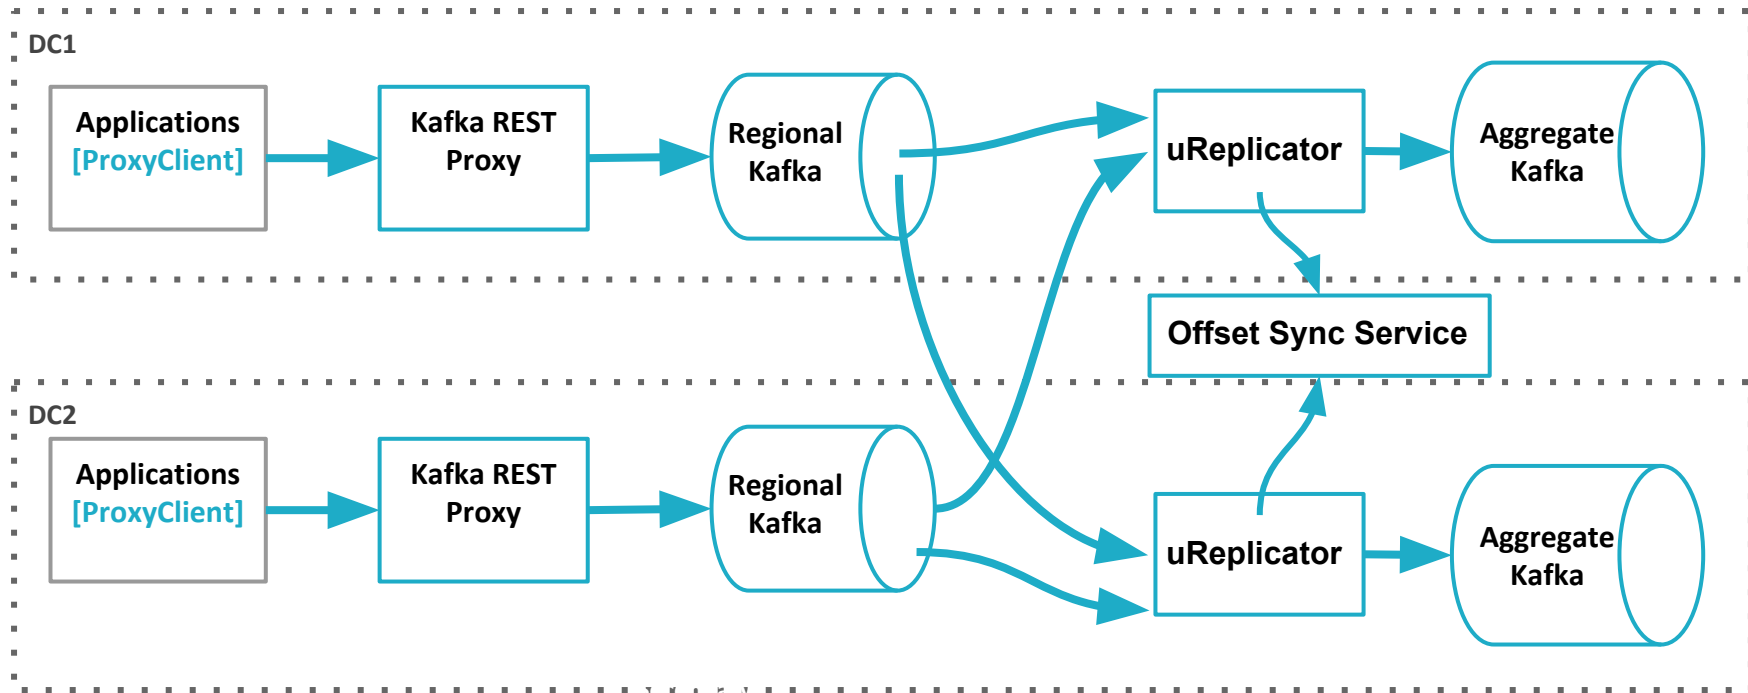
02 Apache Kafka pipeline & replication

03 uReplicator

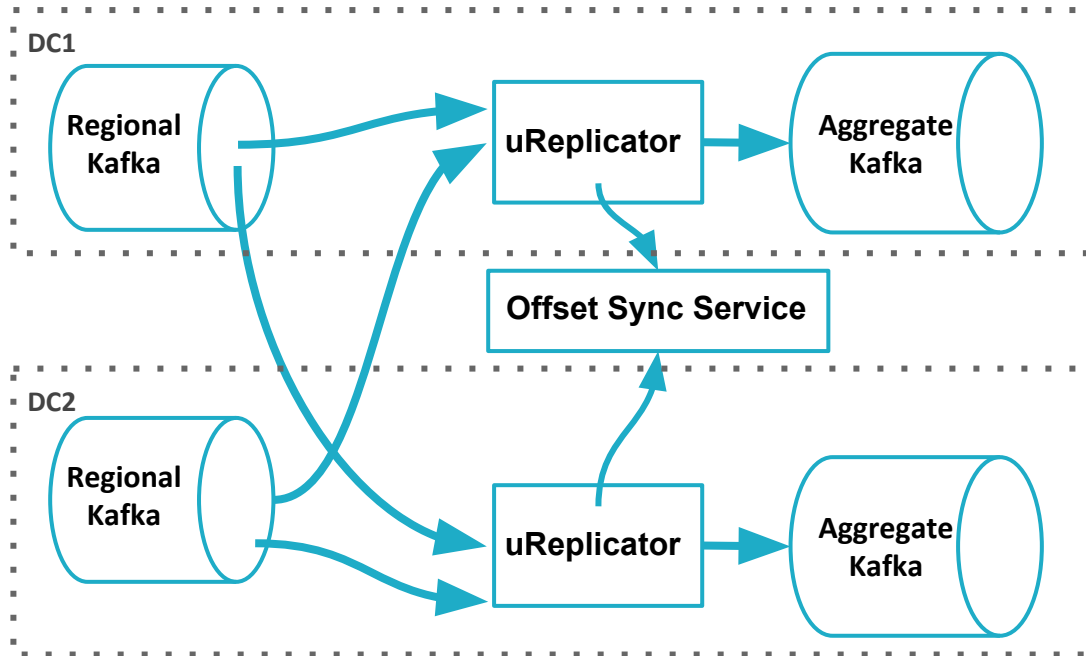
04 Data loss detection

05 Q&A

Apache Kafka Pipeline @ Uber

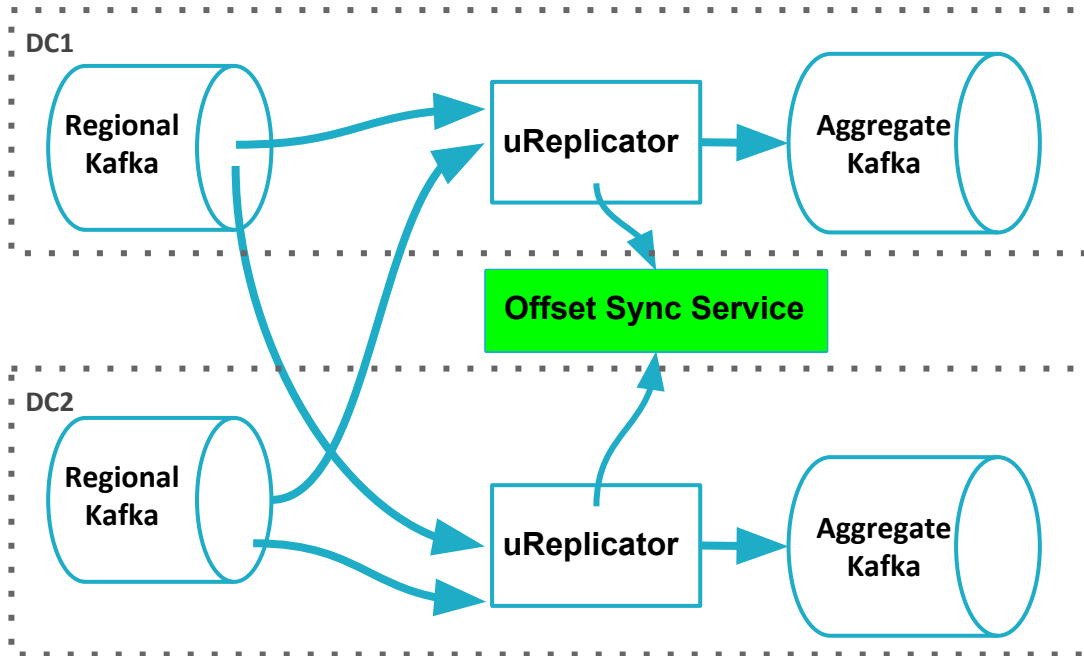


Aggregation



- Global view

Cross-Data Center Failover




During runtime

- uReplicator reports offset mapping to offset sync service
- Offset sync service is all-active and the offset info is replicated across data centers

During failover

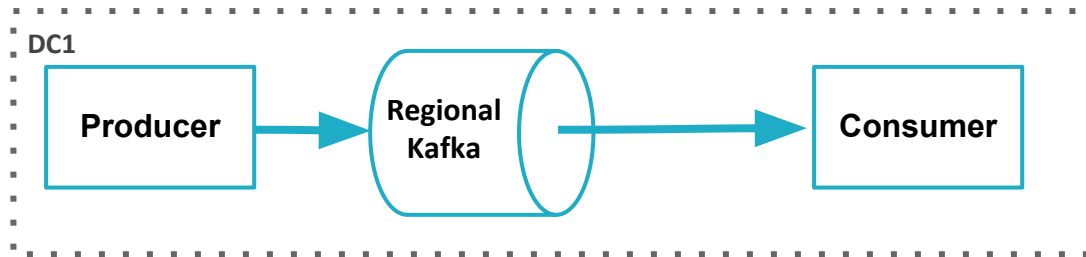
- Consumers ask offset sync service for offsets to resume consumption based on its last commit offsets
- Offset sync service translates offsets between aggregate clusters

Ingestion




- HDFS
- S3

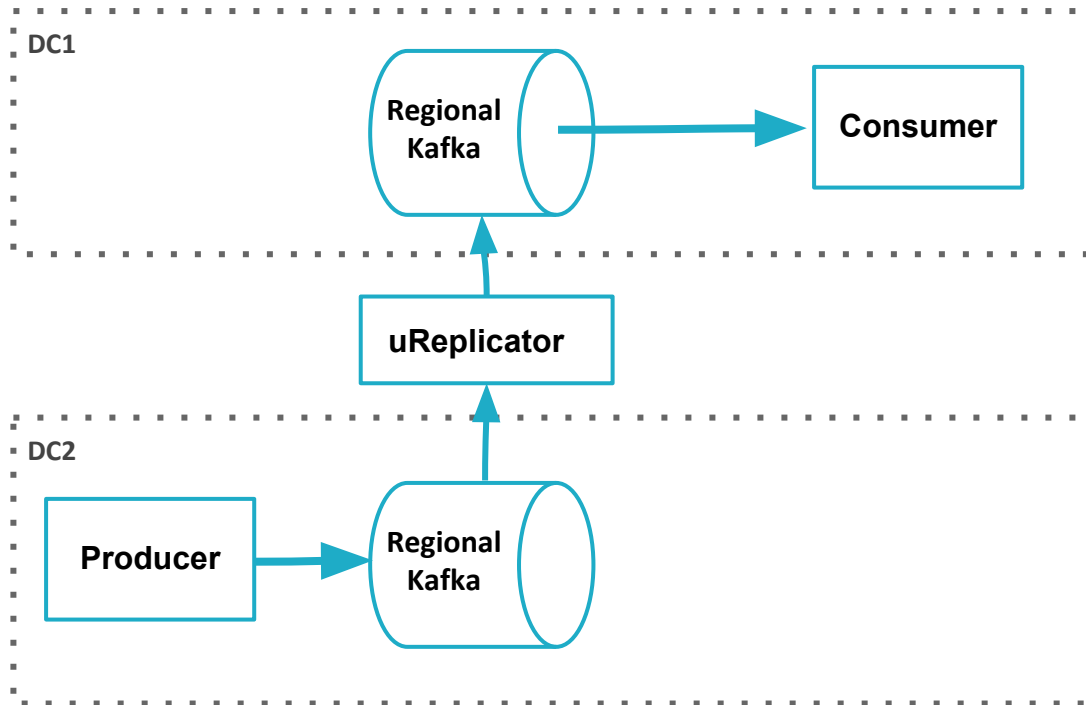
Topic Migration



- Setup uReplicator from new cluster to old cluster
- Move producer
- Move consumer

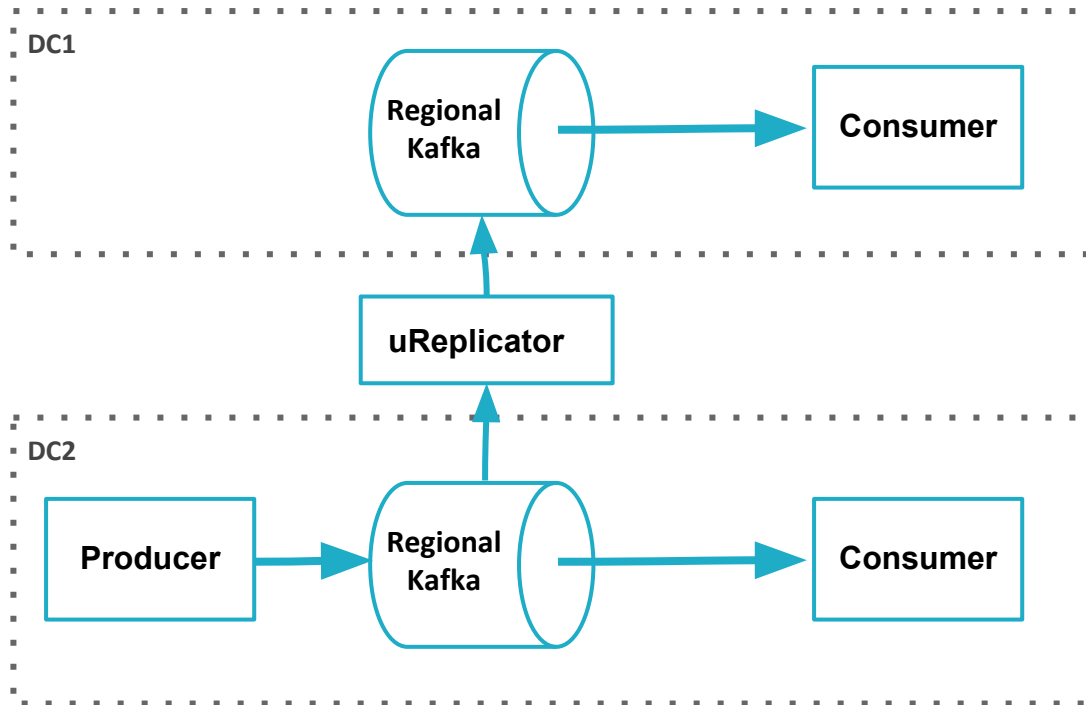


Topic Migration



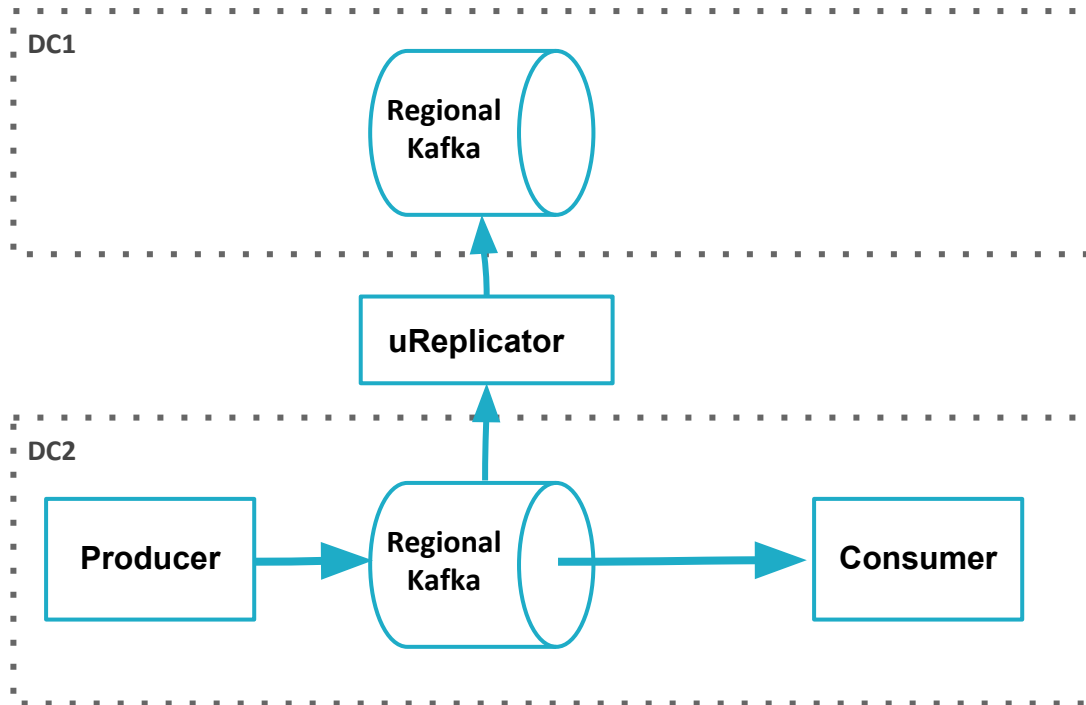
- Setup uReplicator from new cluster to old cluster
- Move producer
- Move consumer
- Remove uReplicator

Topic Migration



- Setup uReplicator from new cluster to old cluster
- Move producer
- Move consumer
- Remove uReplicator

Topic Migration

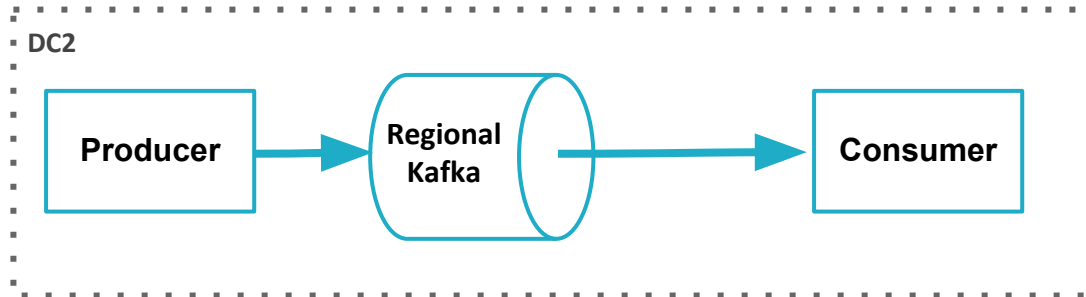


- Setup uReplicator from new cluster to old cluster
- Move producer
- Move consumer
- Remove uReplicator

Topic Migration



- Setup uReplicator from new cluster to old cluster
- Move producer
- Move consumer
- Remove uReplicator



Replication - Use Cases

- Aggregation
 - Global view
 - All-active
 - Ingestion
- Migration
 - Move topics between clusters/DCs

Agenda

01 Apache Kafka at Uber

02 Apache Kafka pipeline & replication

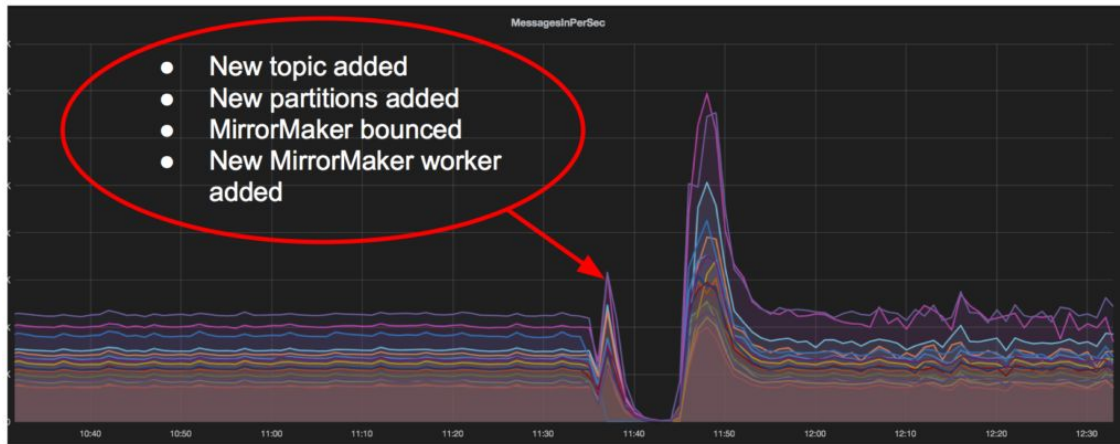
03 uReplicator

04 Data loss detection

05 Q&A

Motivation - MirrorMaker

- Pain point
 - Expensive rebalancing
 - Difficulty adding topics
 - Possible data loss
 - Metadata sync issues



Requirements

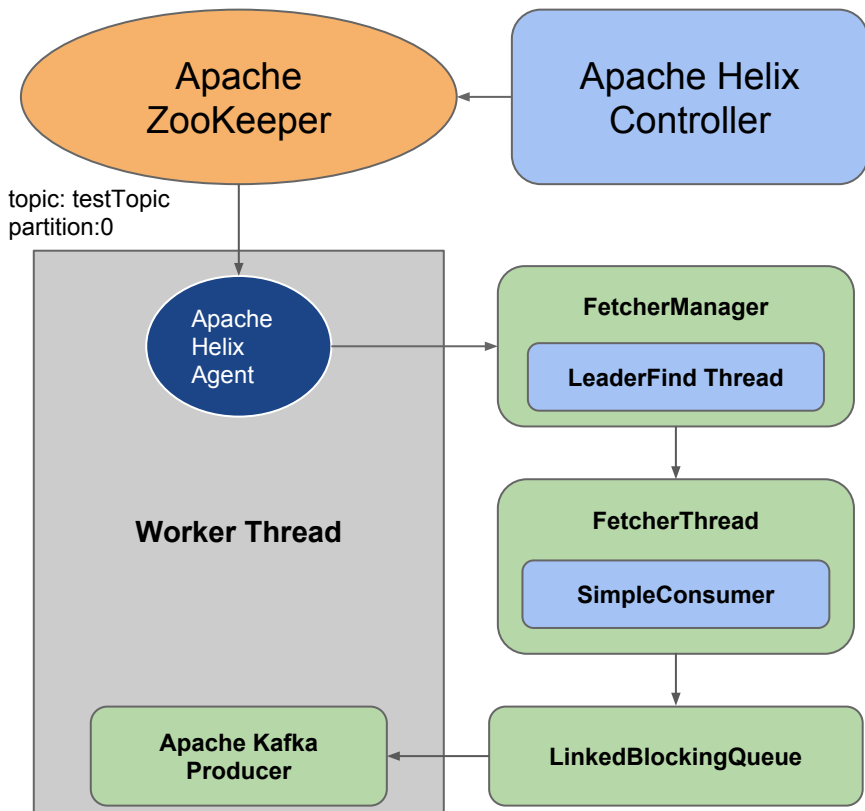
- Stable replication
- Simple operations
- High throughput
- No data loss
- Auditing

Design - uReplicator

- Apache Helix
- uReplicator controller
 - Stable replication
 - Assign topic partitions to each worker process
 - Handle topic/worker changes
 - Simple operations
 - Handle adding/deleting topics

Design - uReplicator

- uReplicator worker
 - Apache Helix agent
 - Dynamic Simple Consumer
 - Apache Kafka producer
 - Commit after flush




Requirements

- ~~Stable replication~~
- ~~Simple operations~~
- High throughput
- ~~No data loss~~
- Auditing



Performance Issues

- Catch-up time is too long (4 hours failover)
 - Full-speed phase: 2~3 hours
 - Long tail phase: 5~6 hours



Problem: Full Speed Phase

- Destination brokers are bound by CPUs




Solution 1: Increase Batch Size


- Destination brokers are bound by CPUs
 - Increase throughput
 - `producer.batch.size`: 64KB => 128KB
 - `producer.linger.ms`: 100 => 1000
- Effect
 - Batch size increases: 10~22KB => 50~90KB
 - Compress rate: 40~62% => 27~35% (compressed size)

Solution 2: 1-1 Partition Mapping

- Round robin
- Topic with N partitions
- N^2 connection
- DoS-like traffic




- Deterministic partition
- Topic with N partitions
- N connection
- Reduce contention




Full Speed Phase Throughput

- First hour: 27.2MB/s => 53.6MB/s per aggregate broker

Before

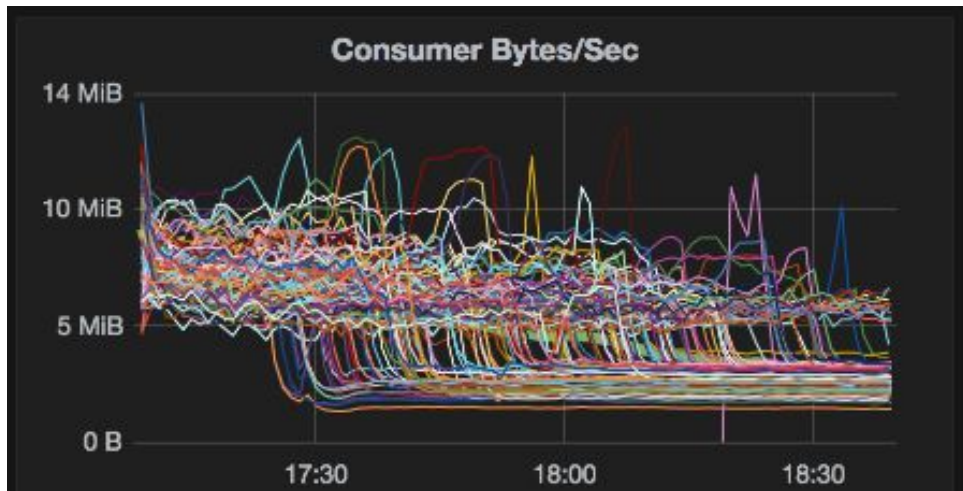


After



Problem 2: Long Tail Phase

- During full speed phase, all workers are busy
- Some workers catch up and become idle, but the others are still busy



Solution 1: Dynamic Workload Balance



- During full speed phase, all workers are busy
- Original partition assignment
 - Number of partitions
 - Heavy partition on the same worker
- Workload-based assignment
 - Total workload when added
 - Source cluster bytes-in-rate
 - Retrieved from Chaperone3
 - Dynamic rebalance periodically
 - Exceeds 1.5 times the average workload

Solution 2: Lag-Feedback Rebalance

- During full speed phase, all workers are busy
- Monitors topic lags
 - Balance the workers based on lags
 - Larger lag = heavier workload
 - Dedicated workers for lagging topics

Dynamic Workload Rebalance

- Periodically adjust workload every 10 minutes
- Multiple lagging topics on the same worker spreads to multiple workers



Requirements

- ~~Stable replication~~
- ~~Simple operations~~
- ~~High throughput~~
- ~~No data loss~~
- Auditing

Requirements

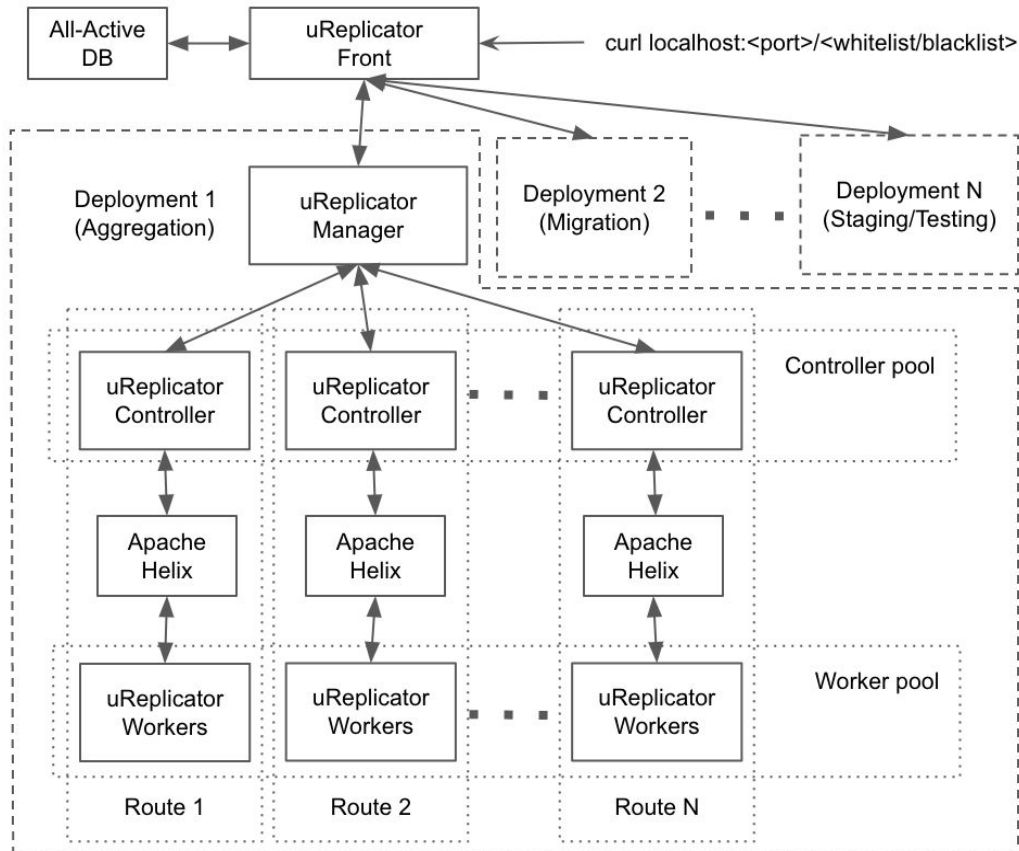
- Stable replication
- Simple operations
- ~~High throughput~~
- ~~No data loss~~
- Auditing

More Issues

- Scalability
 - Number of partitions
 - Up to 1 hour for a full rebalance during rolling restart
- Operation
 - Number of deployments
 - New cluster
- Sanity
 - Loop
 - Double route

Design - Federated uReplicator

- Scalability
 - Multiple route
- Operation
 - Deployment group
 - One manager



Design - Federated uReplicator

- uReplicator Manager
 - Decide when to create new route
 - Decide how many workers in each route
 - Auto-scaling
 - Total workload of route
 - Total lag in the route
 - Total workload / expected workload on worker
 - Automatically add workers to route

Design - Federated uReplicator

- uReplicator Front
 - whitelist/blacklist topics (topic, src, dst)
 - Persist in DB
 - Sanity checks
 - Assigns to uReplicator Manager

Requirements

- ~~Stable replication~~
- ~~Simple operations~~
- ~~High throughput~~
- ~~No data loss~~
- Auditing

Agenda

01 Apache Kafka at Uber

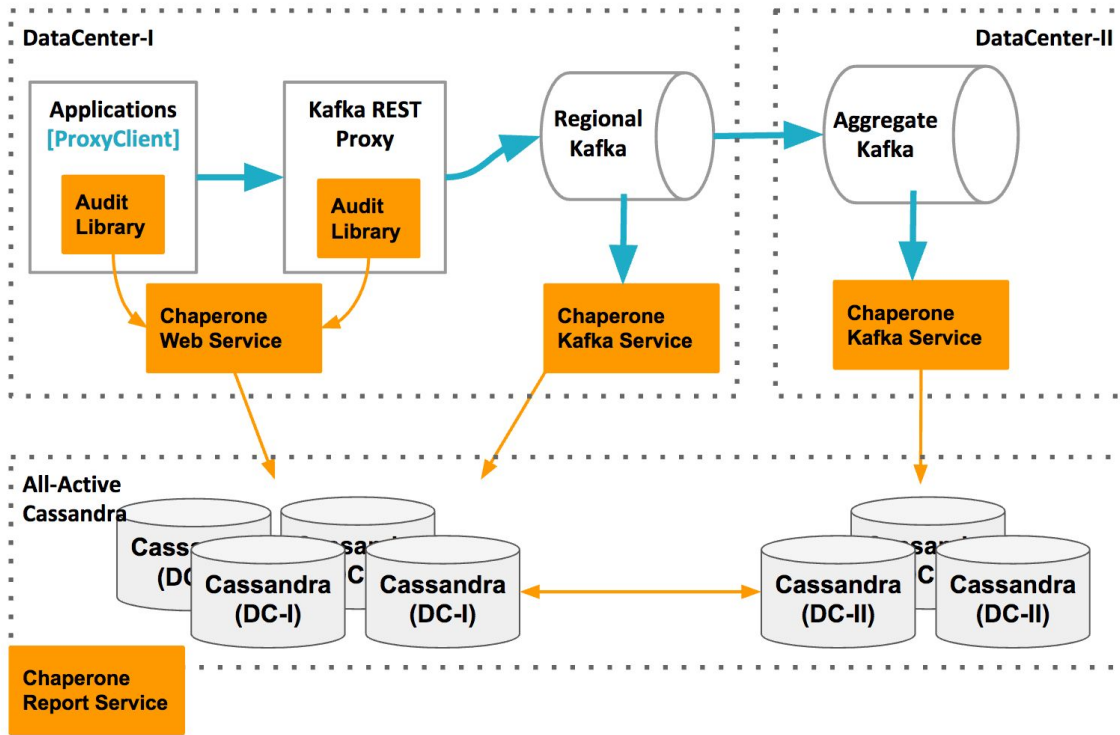
02 Apache Kafka pipeline & replication

03 uReplicator

04 Data loss detection

05 Q&A

Detect data loss



- It checks data as it flows through each tier of the pipeline
- It keeps ts-offset index to support ts-based query
- It also checks latency from the PoV of consumer
- It also provides workload for each topic

Requirements

- ~~Stable replication~~
- ~~Simple operations~~
- ~~High throughput~~
- ~~No data loss~~
- ~~Auditing~~

Blogs & Open Source

- uReplicator
 - Running in production for 2+ years
 - Open sourced: <https://github.com/uber/uReplicator>
 - Blog: <https://eng.uber.com/ureplicator/>
- Chaperone
 - Running in production for 2+ years, audit almost all topics in our clusters
 - Open sourced: <https://github.com/uber/chaperone>
 - Blog: <https://eng.uber.com/chaperone/>

Agenda

01 Apache Kafka at Uber

02 Apache Kafka pipeline & replication

03 uReplicator

04 Data loss detection

05 Q&A

Uber

Proprietary and confidential © 2018 Uber Technologies, Inc. All rights reserved. No part of this document may be reproduced or utilized in any form or by any means, electronic or mechanical, including photocopying, recording, or by any information storage or retrieval systems, without permission in writing from Uber. This document is intended only for the use of the individual or entity to whom it is addressed and contains information that is privileged, confidential or otherwise exempt from disclosure under applicable law. All recipients of this document are notified that the information contained herein includes proprietary and confidential information of Uber, and recipient may not make use of, disseminate, or in any way disclose this document or any of the enclosed information to any person other than employees of addressee to the extent necessary for consultations with authorized personnel of Uber.