Volcano: Bring Batch Capability Into Kubernetes

Da Ma (@k82cn)

Huawei Expert

Geekbang. 极客邦科技 InfoQ

极客邦科技
会议推荐2019

5月

QCon 北京
全球软件开发大会
大会：5月6–8日
培训：5月9–10日

QCon 广州
全球软件开发大会
培训：5月25–26日
大会：5月27–28日

6月

GTLC GLOBAL TECH LEADERSHIP CONFERENCE 上海
技术领导力峰会
时间：6月14–15日

GMTC 北京
全球大前端技术大会
大会：6月20–21日
培训：6月22–23日

7月

ArchSummit 深圳
全球架构师峰会
大会：7月12–13日
培训：7月14–15日

10月

QCon 上海
全球软件开发大会
大会：10月17–19日
培训：10月20–21日

11月

GMTC 深圳
全球大前端技术大会
大会：11月8–9日
培训：11月10–11日

AiCon 北京
全球人工智能与机器学习大会
大会：11月21–22日
培训：11月23–24日

12月

ArchSummit 北京
全球架构师峰会
大会：12月6–7日
培训：12月8–9日

TGO鲲鹏会

汇聚全球科技领导者的高端社群

全球12大城市

850+高端科技领导者

使命 Mission

为社会输送更多优秀的科技领导者

愿景 Vision

构建全球领先的有技术背景优秀人才的学习成长平台

扫描二维码，了解更多内容

About The SPEAKER

Da

Ma

Software Architect

- Kubernetes SIG-Scheduling co-Leader
- Volcano & kube-batch creator
- Expert at Huawei (now)
- Ex-IBM Spectrum CE/L3 Team/Tech Lead
- Jilin University master’s degree, majoring in grid computing and distributed system

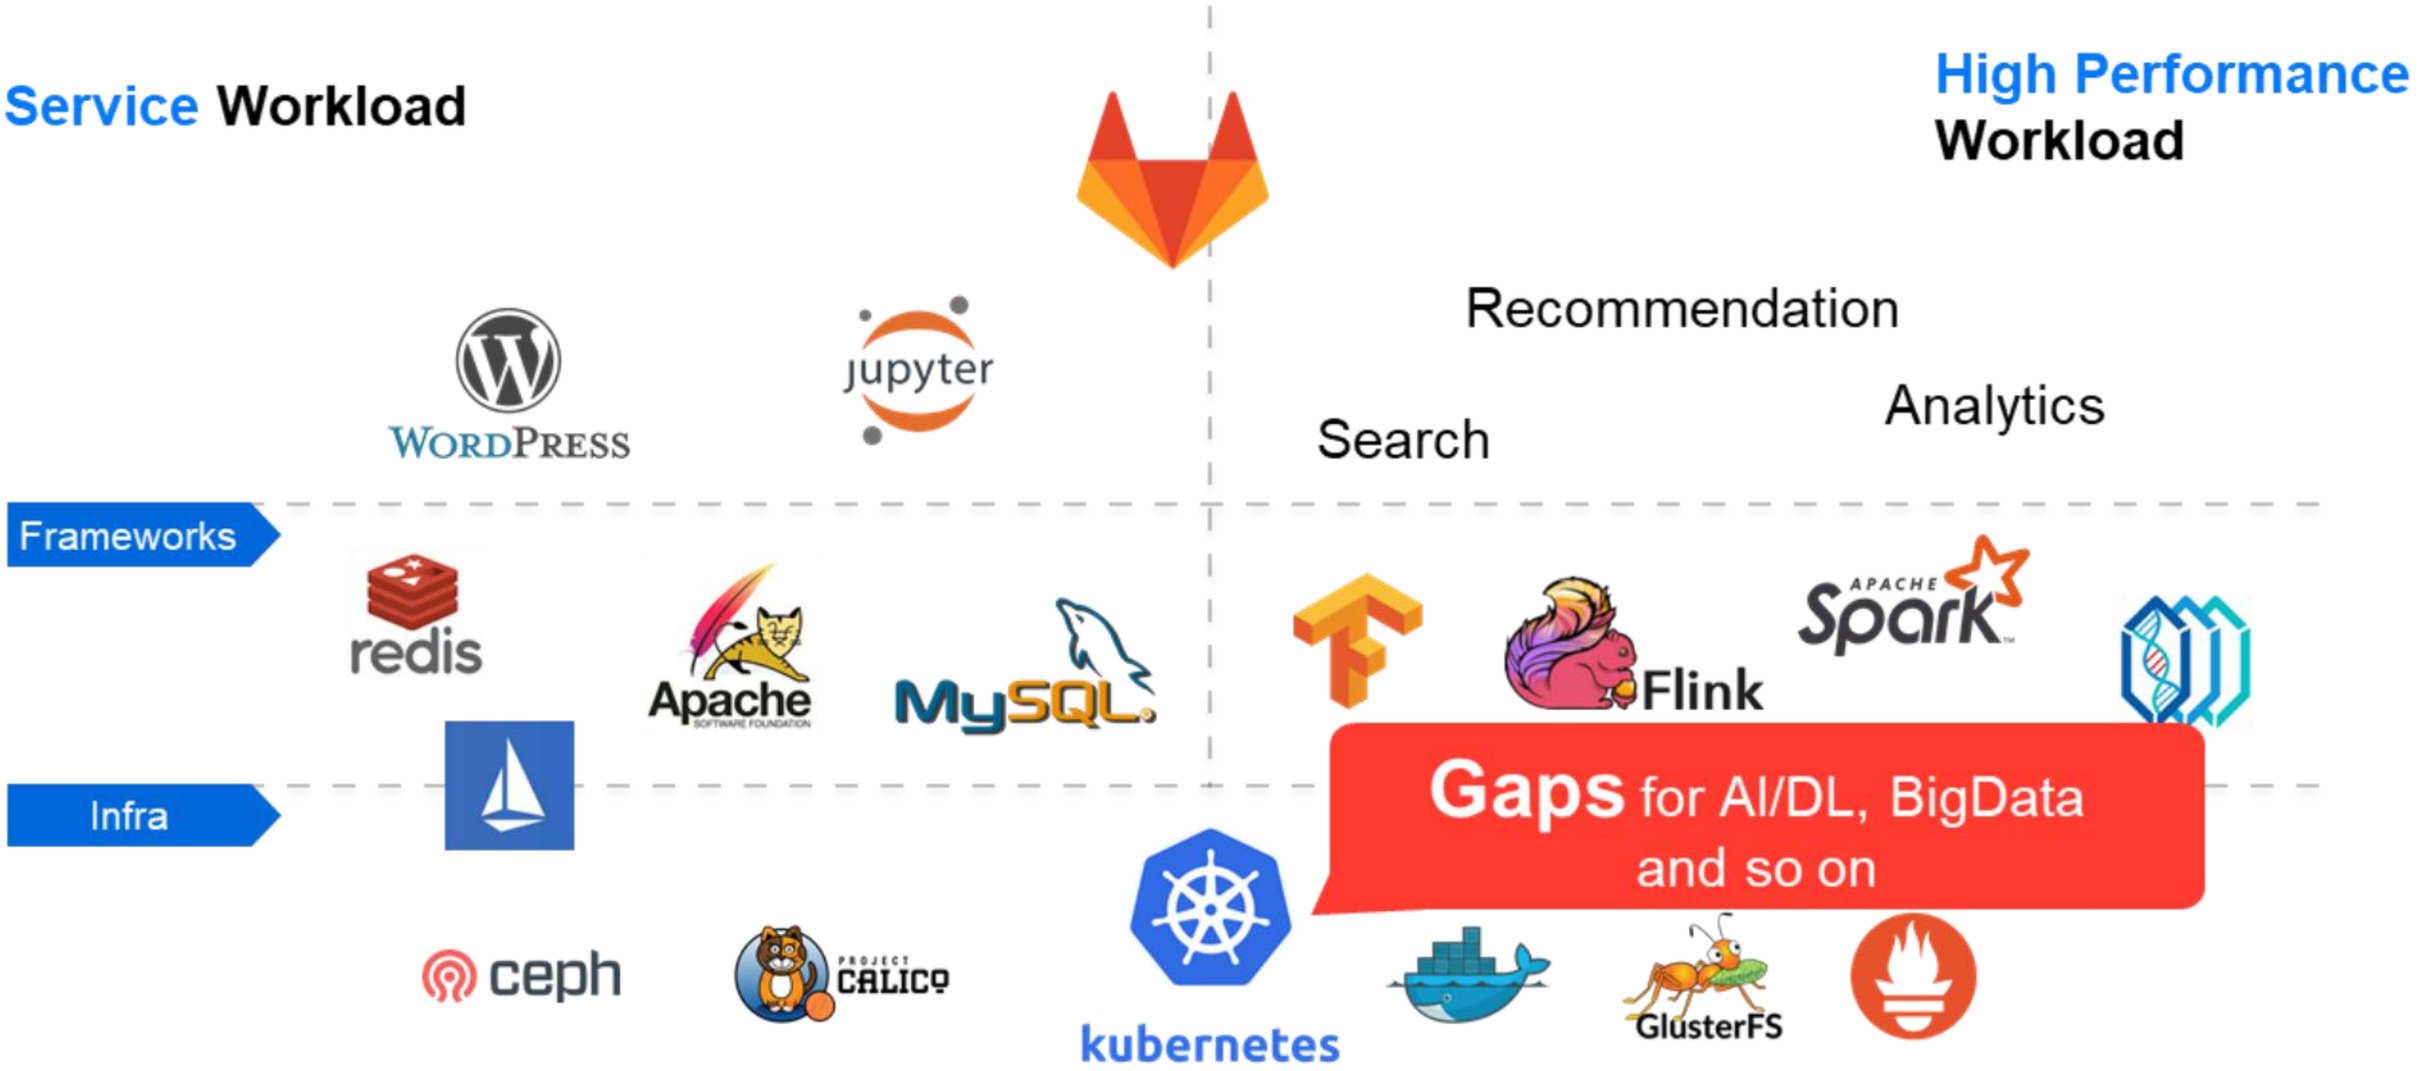
Service Workload
High Performance Workload
Recommendation
Analytics
Search
Frameworks
Infra
WordPress
jupyter
redis
Apache
MySQL
Flink
Spark
ceph
Project Calico
kubernetes
GlusterFS
Gaps for AI/DL, BigData and so on

Gaps

- Job/Queue Management
  - Queue status/configuration
  - Hierarchical queue
  - Job with multiple pod template
  - Lifecycle management of Job, e.g. restart, suspend/resume
  - Error Handling, e.g. restart job if pod failed (MPI, TFJob)
  - Indexed Job
  - Task dependency, e.g. Spark (executor/driver)
  - Delay Pod Creation
  - ...
- Runtime
  - Singularity
  - ...

- Scheduler
  - Coscheduling
  - Faire-share
  - Queue
  - Preemption/Reclaim
  - Reserve/Backfill
  - Topology (network, accelerator)
  - ...
- Others
  - Throughput
  - Round-trip
  - Data locality (Data Aware Scheduling)
  - ...

Volcano: A Kubernetes Native Batch System

Website: https://volcano.sh

Github: http://github.com/volcano-sh/volcano

Twitter: https://twitter.com/volcano_sh

Slack: http://volcano-sh.slack.com

Email: volcano-sh@googlegroups.com

Overall Architecture

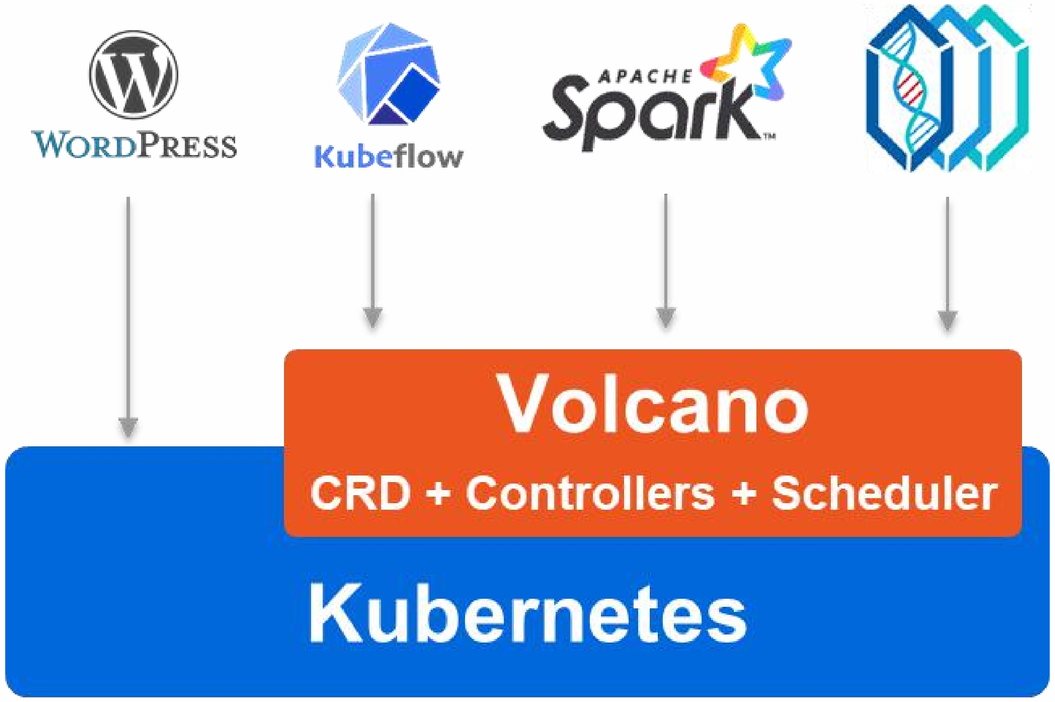

Domain frameworks:

- Deployment/Installation of framework in k8s
- Map framework's terms/concepts into common concept, e.g. Job, Queue
- Enable related features for frameworks, e.g. gang-scheduling for TensorFlow training

Common Service for high performance workload:

- Batch scheduling, e.g. fair-share, gang-scheduling
- Enhanced job management, e.g. multiple pod template, error handling
- Accelerator, e.g. GPU, FPGA
- kubectl plugins, e.g. show Job/Queue information

Overall Architecture

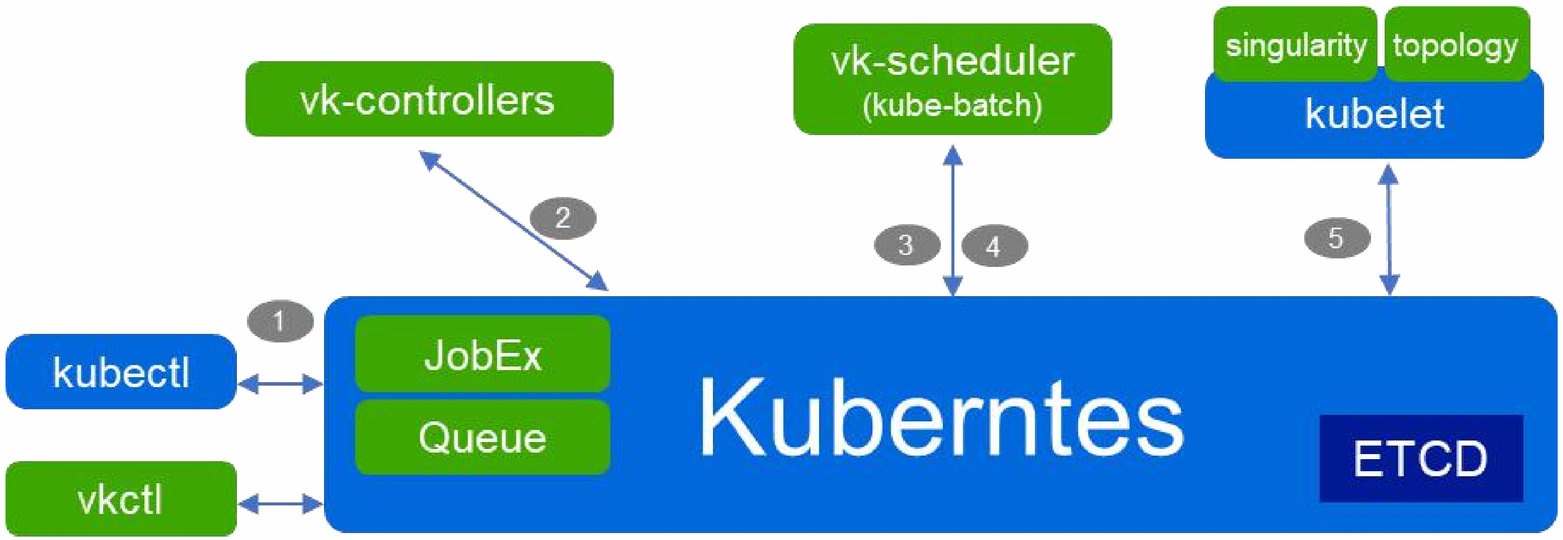

- Kubectl creates a JobEx object in apiserver if all admission passed
- JobExController create Pods based on its replicas and templates
- vk-scheduler get the notification of Pod from apiserver
- vk-scheduler chooses one host for the Pod of JobEx based on its policy
- kubelet gets the notification of Pod from apiserver; and then start the container

- The policy in vk-scheduler is pluggable, e.g. DRF, Priority, Gang
- vk-controllers includes JobExController, QueueController
- Volcano handles high performance workload

Scenarios: MPI

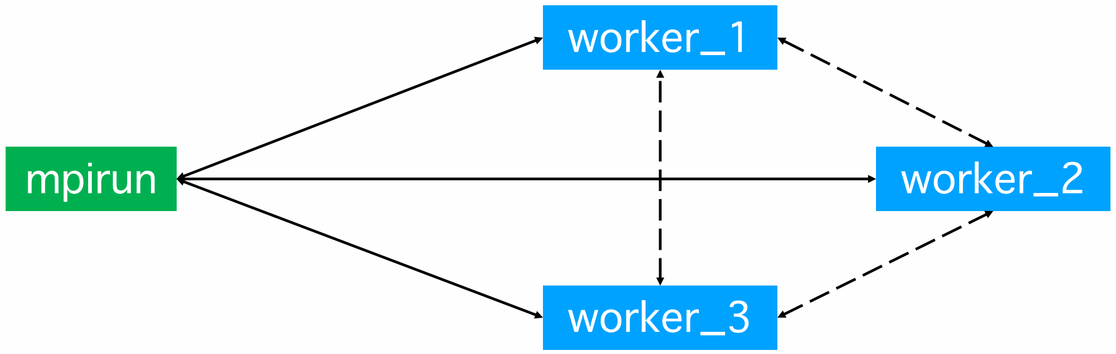

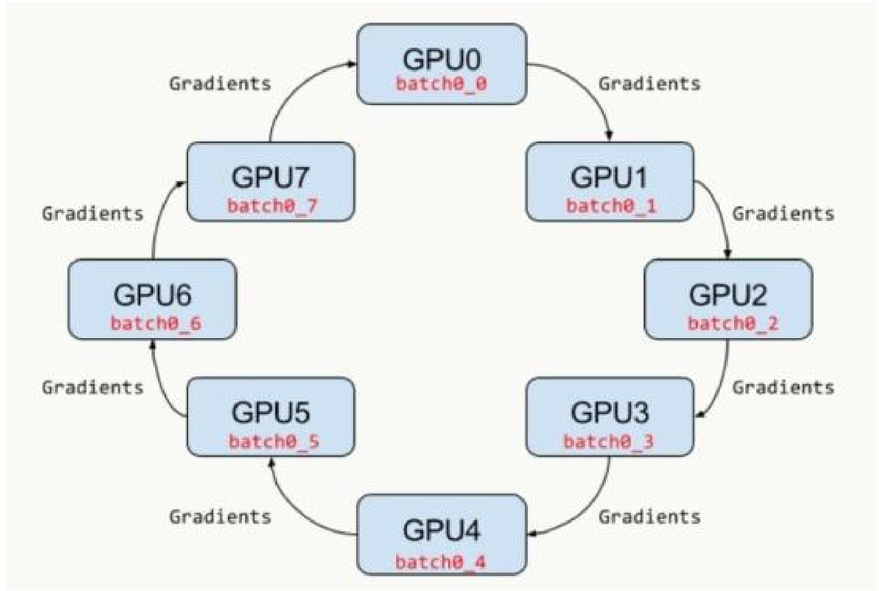

- Multiple Pod Template
- Lifecycle Policy
- Gang-scheduling

- ssh or kubectl
- Complete job when mpirun completed
- Headless service

Scenarios: MPI

```yaml
apiVersion: batch.volcano.sh/v1alpha1
kind: Job
metadata:
  name: lm-mpi-job
  labels:
    # 根据业务需要设置作业类型
    "volcano.sh/job-type": "MPI"
spec:
  # 设置最小需要的服务（小于总replicas数）
  minAvailable: 3
  schedulerName: volcano
  plugins:
    # 提供 ssh 免密认证
    ssh: []
    # 提供运行作业所需要的网络信息，hosts文件，headless service等
    svc: []
  # 如果有pod被 杀死，重启整个作业
  policies:
    - event: PodEvicted
      action: RestartJob
  tasks:
    - replicas: 1
      name: mpimaster
      # 当 mpiexec 结束，认识整个mpi作业结束
      policies:
        - event: TaskCompleted
          action: CompleteJob
      template:
        spec:
          # Volcano 的信息会统一放到 /etc/volcano 目录下
          containers:
            - command:
                - /bin/sh
                - -c
                - |
```

```
Pods:
------------------------------
NAME                                        READY   STATUS      RESTARTS   AGE
lm-mpi-job-mpimaster-0                      0/1     Completed   3          2m
spark-operator-sparkoperator-f78854b64-rh52d   1/1     Running     0          1d


Volcano Jobs:
------------------------------
Name         Creation              Phase       JobType     Replicas    Min   Pending   Running   Succeeded
lm-mpi-job   2019-06-19 20:55:33   Completed   MPI         3           3     0         0         1


m00483107@M00483107 MINGW64 /d/workspace/src/volcano.sh/volcano/docs/samples/kubecon-2019-china/mpi-sample (kub
$ kc logs lm-mpi-job-mpimaster-0
Warning: Permanently added 'lm-mpi-job-mpiworker-0.lm-mpi-job,172.16.0.22' (ECDSA) to the list of known hosts.
Warning: Permanently added 'lm-mpi-job-mpiworker-1.lm-mpi-job,172.16.0.46' (ECDSA) to the list of known hosts.
Hello world from processor lm-mpi-job-mpiworker-0, rank 0 out of 2 processors
Hello world from processor lm-mpi-job-mpiworker-1, rank 1 out of 2 processors
```

- The workers are deleted by job controller because of sshd
- The pod of mpiexec/mpirun will not be deleted for output
- The pod of mpiexec/mpirun may restart few times because of network setup delay

```
m00483107@M00483107 MINGW64 /d/workspace/src/volcano.sh/volcano/docs/samples/kubecon-2019-china/mpi-sample (kubecon-2019-china-demo)
$ watch.sh
Nodes:


-------------------------------------------------------------------------------------------------------
               | 192.168.23.24          | 192.168.51.67          | 192.168.57.163         | 192.168.63.50          |
-------------------------------------------------------------------------------------------------------
 Alloctable    | cpu: 3920m             | cpu: 3920m             | cpu: 3920m             | cpu: 3920m             |
               | mem: 6265300Ki         | mem: 6265300Ki         | mem: 6265300Ki         | mem: 6265300Ki         |
-------------------------------------------------------------------------------------------------------
 Idle          | cpu: 2420m             | cpu: 1920m             | cpu: 3420m             | cpu: 2920m             |
               | mem: 3643860Ki         | mem: 4168148Ki         | mem: 5741012Ki         | mem: 4168148Ki         |
-------------------------------------------------------------------------------------------------------


Pods:
------------------------------
NAME                                     READY   STATUS    RESTARTS   AGE
spark-operator-sparkoperator-f78854b64-rh52d   1/1     Running   0          1d


Volcano Jobs:
------------------------------
No resources found

m00483107@M00483107 MINGW64 /d/workspace/src/volcano.sh/volcano/docs/samples/kubecon-2019-china/mpi-sample (kubecon-2019-china-demo)
$
```

Scenarios: Faire Share

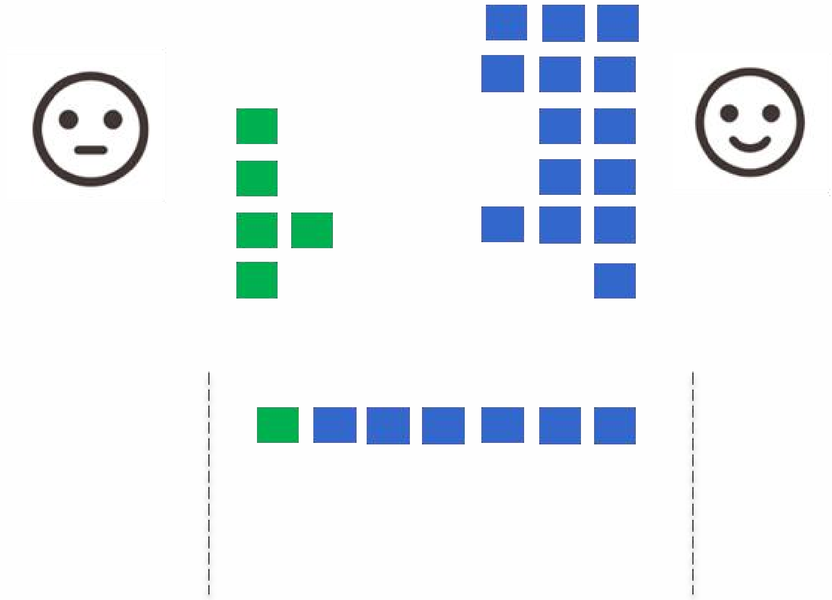

The more workload, the more resources???

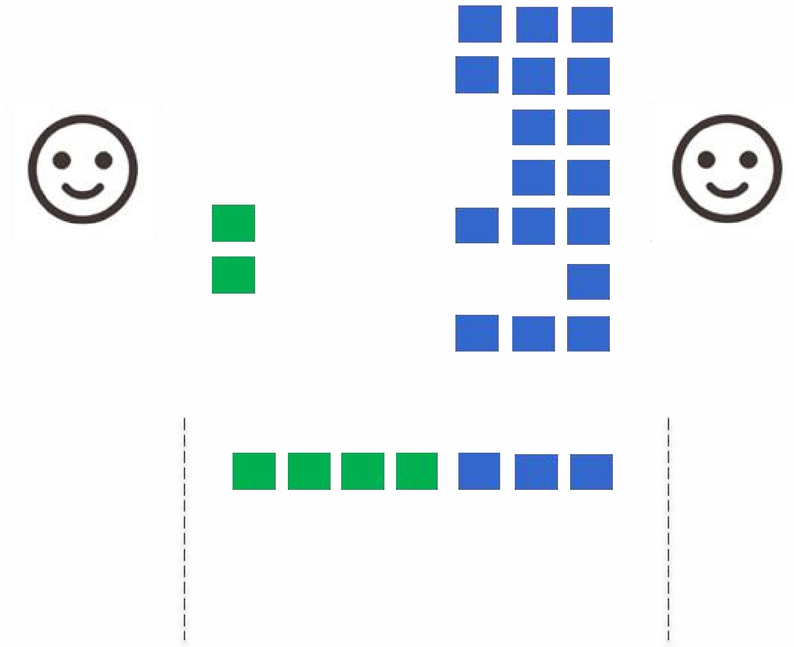

Share resources by weight !!!

```
m00483107@M00483107 MINGW64 /d/workspace/src/volcano.sh/volcano/docs/samples/kubecon-2019-china/drf (kubecon-2019-china-demo)
$ watch.sh
Nodes:

-------------------------------------------------------------------------------------------------------
              | 192.168.23.24         | 192.168.51.67         | 192.168.57.163        | 192.168.63.50         |
-------------------------------------------------------------------------------------------------------
 Alloctable   | cpu: 3920m            | cpu: 3920m            | cpu: 3920m            | cpu: 3920m            |
              | mem: 6265300Ki        | mem: 6265300Ki        | mem: 6265300Ki        | mem: 6265300Ki        |
-------------------------------------------------------------------------------------------------------
 Idle         | cpu: 2420m            | cpu: 1920m            | cpu: 3420m            | cpu: 2920m            |
              | mem: 3643860Ki        | mem: 4168148Ki        | mem: 5741012Ki        | mem: 4168148Ki        |
-------------------------------------------------------------------------------------------------------


Pods:
-------------------------------
NAME                                     READY   STATUS      RESTARTS   AGE
lm-horovod-job-master-0                  0/1     Completed   1          41m
lm-mpi-job-mpimaster-0                   0/1     Completed   3          3m
spark-operator-sparkoperator-f78854b64-rh52d   1/1     Running     0          1d
spark-pi-2-driver                        0/1     Completed   0          7m
spark-pi-3-driver                        0/1     Completed   0          7m
spark-pi-driver                          0/1     Completed   0          7m


Volcano Jobs:
-------------------------------
Name            Creation                Phase       JobType     Replicas    Min    Pending    Running    Succeeded    Failed    RetryCount
lm-horovod-job    2019-06-19 21:11:04     Completed   Horovod     4           4      0          0          1            0         2

lm-mpi-job        2019-06-19 21:53:39     Completed   MPI         4           4      0          0          1            0         0


m00483107@M00483107 MINGW64 /d/workspace/src/volcano.sh/volcano/docs/samples/kubecon-2019-china/drf (kubecon-2019-china-demo)
$
```

Gang-scheduling: Job Execution Time

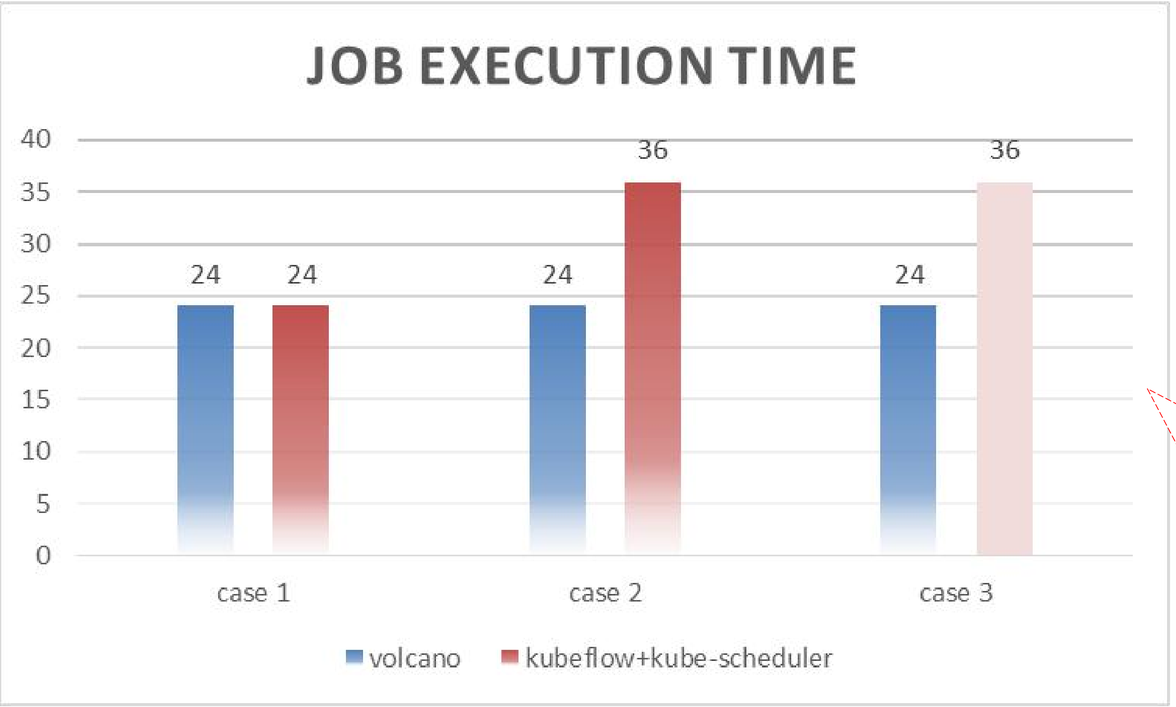

- Case 1: 1 job with 2ps + 4workers
- **Case 2: 2 jobs with 2ps + 4workers**
- **Case 3: 5 jobs with 2ps + 4workers**

- No enough resource for 2 Jobs to run concurrently; one of them **wasting** resources without Gang-Scheduling !
- 2 of 5 jobs was finished because of deadlock (+20 hours)

http://status.openlabtesting.org/builds?project=theopenlab%2Fvolcano

Task-Topology + Binpack

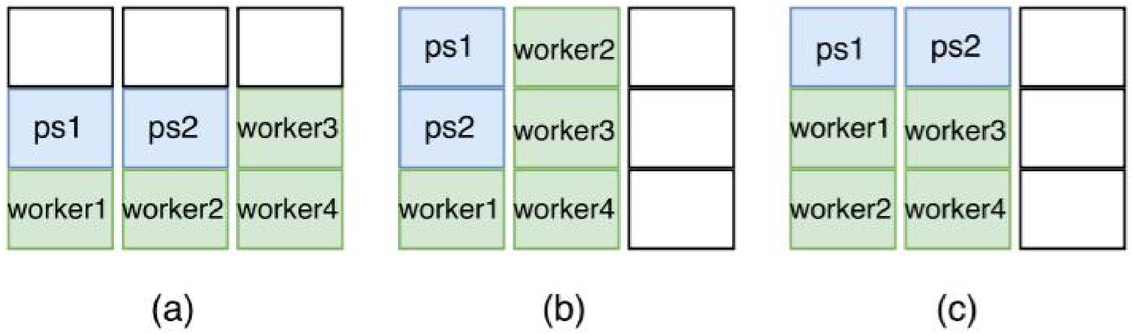

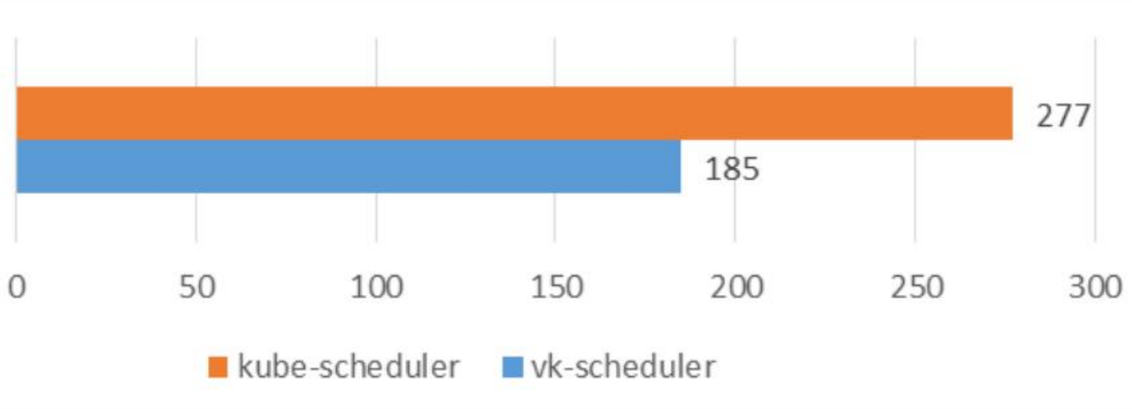

- The execution time of 3 jobs in total; 2ps + 4workers for each job
- The execution time is unstable when tested by default scheduler
- The improvement dependent on data exchanges between pods
- Task-topology within a Job also improved scheduler's performance
- Open Source at [volcano-sh/volcano#272]

Reference: "Optimus: An Efficient Dynamic Resource Scheduler for Deep Learning Clusters"

Integrations

Framework	Status	API
MPI	Done	Volcano Job
Horovod	Done	Volcano Job
Kubeflow/tf-operator	Done	PodGroup
Kubeflow/arena	Done	Volcano Job
Spark-Operator	On-going	PodGroup
Cromwell	On-going	Volcano Job
PaddlePaddle	On-going	Volcano Job
...	On-going	Volcano Job / PodGroup

Pipeline

- GPU Share/Topology
- Job Management
- Queue Management
- Hierarchical Queue
- Preemption/Reclaim
- ……

Call for Contribution

极客时间全部课程任学

喊老板来买单！

- 精选 13+ 热门职位的学习路径，包括架构、运维、前端工程师等
- 根据不同技术岗位能力模型匹配合适的课程
- 一键设置购买条件，成员按需选课，自主制定学习计划
- 享充值满赠优惠，帮老板省钱，团队免费学习

立即申请

THANKS

Geekbang. 极客邦科技 InfoQ